

Import Packages

```
In [1]: import pandas as pd
import sklearn
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression, LogisticRegression
from sklearn.metrics import mean_squared_error, accuracy_score
from sklearn.metrics import accuracy_score, classification_report
from sklearn.ensemble import RandomForestClassifier
```

```
In [2]: data = pd.read_csv('finalETLdata.csv')
data
```

Out[2]:

Skillset4	Skillset5	Location	...	Language1	Language2	Language3	Language4	Language5	ResearchInterests	Attendance	SkillCount	LanguageCount
known	Unknown	New York	...	Chinese	Japanese	Spanish	German	French	Biomedical Engineering	0.53	1	5
Data analysis	Programming	Boston	...	French	English	Chinese	Spanish	Japanese	Urban Planning	0.37	5	5
Artistic	Problem Solving	Chicago	...	Spanish	Japanese	German	French	Unknown	Nanotechnology	0.25	5	4
known	Unknown	Chicago	...	Japanese	Chinese	Spanish	French	Unknown	Space Exploration	0.14	3	4
known	Unknown	Chicago	...	English	Unknown	Unknown	Unknown	Unknown	Climate Change	0.15	1	1
...
Public speaking	Data Analysis	Chicago	...	French	Chinese	English	Unknown	Unknown	Astrophysics	0.07	5	3
known	Unknown	New York	...	Spanish	German	Chinese	Unknown	Unknown	Data Science	0.79	2	3
Leadership	Problem Solving	New York	...	Spanish	German	Japanese	Chinese	French	Space Exploration	0.29	5	5
Leadership	Problem Solving	Houston	...	German	Unknown	Unknown	Unknown	Unknown	Space Exploration	0.38	5	1
known	Unknown	New York	...	Japanese	Spanish	Chinese	French	German	Bioinformatics	0.03	1	5

```
In [3]: df = pd.DataFrame(data)
df
```

Out [3]:

	StudentID	Name	AcademicInterest	ExtracurricularActivities	Skillset1	Skillset2	Skillset3	Skillset4	Skillset5	Location	...	LanguageCount	
	0	1	Aaron Jordan	Psychology	Debate Club	Problem Solving	Unknown	Unknown	Unknown	Unknown	New York	...	5
	1	2	Kevin Smith	Psychology	Debate Club	Leadership	Problem Solving	Public Speaking	Data Analysis	Programming	Boston	...	5
	2	3	William May	History	Volunteer Group	Data Analysis	Leadership	Public Speaking	Artistic	Problem Solving	Chicago	...	4
	3	4	Valerie Miranda	Computer Science	Volunteer Group	Public Speaking	Data Analysis	Problem Solving	Unknown	Unknown	Chicago	...	3
	4	5	Samantha Hill	Computer Science	Sports Team	Data Analysis	Unknown	Unknown	Unknown	Unknown	Chicago	...	5

	995	996	Justin Jordan	History	Sports Team	Leadership	Artistic	Programming	Public Speaking	Data Analysis	Chicago	...	5
	996	997	Lauren Baker	Psychology	Art Club	Data Analysis	Leadership	Unknown	Unknown	Unknown	New York	...	5
	997	998	Charles Lewis	Mathematics	Volunteer Group	Public Speaking	Programming	Artistic	Leadership	Problem Solving	New York	...	5
	998	999	Christopher Lewis	Physics	Sports Team	Public Speaking	Data Analysis	Artistic	Leadership	Problem Solving	Houston	...	5
	999	1000	Corey Bond	History	Coding Club	Public Speaking	Unknown	Unknown	Unknown	Unknown	New York	...	5

1000 rows x 23 columns

Label Encoding

```
In [4]: from sklearn.preprocessing import LabelEncoder

# Define the columns you want to apply label encoding to
columns_to_encode = ['AcademicInterest', 'ExtracurricularActivities', 'YearOfStudy', 'Major', 'ResearchInterests']

# Initialize LabelEncoder
label_encoder = LabelEncoder()

# Apply label encoding to selected columns
for column in columns_to_encode:
    df[column] = label_encoder.fit_transform(df[column])

# Display the modified DataFrame
df
```

```
Out[4]:
```

killset5	Location	...	Language1	Language2	Language3	Language4	Language5	ResearchInterests	Attendance	SkillCount	LanguageCount	CompositeS
nknown	New York	...	Chinese	Japanese	Spanish	German	French	4	0.53	1	5	
amming	Boston	...	French	English	Chinese	Spanish	Japanese	24	0.37	5	5	
Problem Solving	Chicago	...	Spanish	Japanese	German	French	Unknown	15	0.25	5	4	
nknown	Chicago	...	Japanese	Chinese	Spanish	French	Unknown	22	0.14	3	4	
nknown	Chicago	...	English	Unknown	Unknown	Unknown	Unknown	6	0.15	1	1	
...	
Data Analysis	Chicago	...	French	Chinese	English	Unknown	Unknown	1	0.07	5	3	
nknown	New York	...	Spanish	German	Chinese	Unknown	Unknown	9	0.79	2	3	
Problem Solving	New York	...	Spanish	German	Japanese	Chinese	French	22	0.29	5	5	
Problem Solving	Houston	...	German	Unknown	Unknown	Unknown	Unknown	22	0.38	5	1	
nknown	New York	...	Japanese	Spanish	Chinese	French	German	3	0.03	1	5	

Task 1: GPA prediction

```
In [5]: # Define predictor variables (X) and target variable (y) for GPA prediction
X_gpa = df[['StudentID', 'SkillCount', 'LanguageCount', 'CompositeScore']]
y_gpa = df['GPA']

In [6]: # Split the data into training and testing sets for GPA prediction
X_train_gpa, X_test_gpa, y_train_gpa, y_test_gpa = train_test_split(X_gpa, y_gpa, test_size=0.2, random_state=42)

In [7]: # Train linear regression model for GPA prediction
lr_gpa = LinearRegression()
lr_gpa.fit(X_train_gpa, y_train_gpa)

# Evaluate linear regression model for GPA prediction
y_pred_gpa = lr_gpa.predict(X_test_gpa)
mse_gpa = mean_squared_error(y_test_gpa, y_pred_gpa)
print("Mean Squared Error (GPA Prediction):", mse_gpa)
```

Mean Squared Error (GPA Prediction): 0.006415792969882901

```
In [8]: # Step 5: Make Predictions
# Use the trained model to predict the next GPA of each student
next_gpa_predictions = lr_gpa.predict(X_gpa)
df['NextGPA'] = next_gpa_predictions

# Display the DataFrame with predicted next GPA
print(df[['StudentID', 'GPA', 'Attendance', 'NextGPA']])
```

	StudentID	GPA	Attendance	NextGPA
0	1	3.27	0.53	3.289579
1	2	3.17	0.37	3.118852
2	3	2.09	0.25	2.324098
3	4	2.56	0.14	2.561646
4	5	2.01	0.15	2.201725
..
995	996	2.21	0.07	2.270511
996	997	3.62	0.79	3.685658
997	998	3.09	0.29	2.999678
998	999	3.17	0.38	3.116268
999	1000	2.16	0.03	2.217429

[1000 rows x 4 columns]

Task 2: Pass Fail Prediction

```
In [9]: # Define predictor variables (X) and target variable (y) for pass/fail prediction
threshold = 3.0 # Threshold GPA value for pass/fail classification
passOrFail = (df['GPA'] >= threshold).astype(int) # Create binary target variable for pass/fail classification

# Exclude GPA and PassOrFail from predictor variables
X_train_pf, y_train_pf = train_test_split(X, y, test_size=0.2, random_state=42)
X_test_pf, y_test_pf = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [10]: # Split the data into training and testing sets for pass/fail prediction
X_train_pf, X_test_pf, y_train_pf, y_test_pf = train_test_split(X_pf, y_pf, test_size=0.2, random_state=42)
```

```
In [11]: # Train logistic regression model for pass/fail prediction
lr_pf = LogisticRegression()
lr_pf.fit(X_train_pf, y_train_pf)

# Step 4: Evaluate Model Performance
# For pass/fail classification
y_pred_pf = lr_pf.predict(X_test_pf)
accuracy_pf = accuracy_score(y_test_pf, y_pred_pf)
print("Accuracy (Pass/Fail Classification):", accuracy_pf)
print("Classification Report (Pass/Fail Classification):\n", classification_report(y_test_pf, y_pred_pf))
```

```
Accuracy (Pass/Fail Classification): 0.985
Classification Report (Pass/Fail Classification):
              precision    recall  f1-score   support

    0           0.98         0.99         0.98         91
    1           0.99         0.98         0.99        109

 accuracy               0.98         0.98         0.98        200
 macro avg              0.98         0.99         0.98        200
 weighted avg           0.99         0.98         0.99        200
```

```
In [12]: # Display the DataFrame with pass or fail predictions
print(df[['StudentID', 'GPA', 'Attendance', 'CompositeScore', 'NextGPA', 'PassOrFail']])
```

	StudentID	GPA	Attendance	CompositeScore	NextGPA	PassOrFail
0	1	3.27	0.53	1.90	3.289579	1
1	2	3.17	0.37	1.77	3.118852	1
2	3	2.09	0.25	1.17	2.324098	0
3	4	2.56	0.14	1.35	2.561646	0
4	5	2.01	0.15	1.08	2.201725	0
..
995	996	2.21	0.07	1.14	2.270511	0
996	997	3.62	0.79	2.21	3.685658	1
997	998	3.09	0.29	1.69	2.999678	1
998	999	3.17	0.38	1.78	3.116268	1
999	1000	2.16	0.03	1.10	2.217429	0

[1000 rows x 6 columns]

Task 3: Engagement Predictions

```
In [13]: # Filter out students who have passed based on GPA
engaged_students = df[df['GPA'] >= threshold]
```

```
In [14]: # Define engagement categories based on specified criteria
def define_engagement(row):
    if row['ExtracurricularActivities'] > 0 and row['ResearchInterests'] is not None:
        if row['SkillCount'] > 3 and row['LanguageCount'] > 3:
            return "Extremely Engaged"
        elif row['SkillCount'] > 1 and row['LanguageCount'] > 1:
            return "Moderately Engaged"
        elif row['SkillCount'] == 1 and row['LanguageCount'] == 1:
            return "Low Engaged"
        else:
            return "Not Engaged"
    else:
        return "Not Engaged"

# Apply the define_engagement function to create the 'Engagement' variable
engaged_students['Engagement'] = engaged_students.apply(lambda row: define_engagement(row), axis=1)

# Label encode the 'Engagement' variable
label_encoder = LabelEncoder()
engaged_students['Engagement'] = label_encoder.fit_transform(engaged_students['Engagement'])

engaged_students
```

/var/folders/33/03_543s9719fh3y8mrxhst580000gn/T/ipykernel_46941/1561929157.py:16: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
engaged_students['Engagement'] = engaged_students.apply(lambda row: define_engagement(row), axis=1)
```

/var/folders/33/03_543s9719fh3y8mrxhst580000gn/T/ipykernel_46941/1561929157.py:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
engaged_students['Engagement'] = label_encoder.fit_transform(engaged_students['Engagement'])
```

Out[14]:

	StudentID	Name	AcademicInterest	ExtracurricularActivities	Skillset1	Skillset2	Skillset3	Skillset4	Skillset5	Location	...	Langu
0	1	Aaron Jordan	5	2	Problem Solving	Unknown	Unknown	Unknown	Unknown	New York	...	Gi
1	2	Kevin Smith	5	2	Leadership	Problem Solving	Public Speaking	Data Analysis	Programming	Boston	...	Sf
6	7	Jason Hale	2	0	Public Speaking	Problem Solving	Data Analysis	Artistic	Unknown	New York	...	Gi
7	8	Anthony Wright	1	5	Programming	Public Speaking	Artistic	Leadership	Unknown	Houston	...	Unf
9	10	Peter Yu	0	4	Programming	Data Analysis	Problem Solving	Leadership	Unknown	Los Angeles	...	Unf
...
990	991	Gabriel Nolan	5	4	Programming	Problem Solving	Unknown	Unknown	Unknown	Los Angeles	...	Unf
993	994	Jeffrey Williams	1	4	Data Analysis	Artistic	Unknown	Unknown	Unknown	New York	...	Unf
996	997	Lauren Baker	5	0	Data Analysis	Leadership	Unknown	Unknown	Unknown	New York	...	Unf
997	998	Charles Lewis	3	5	Public Speaking	Programming	Artistic	Leadership	Problem Solving	New York	...	Cl
998	999	Christopher Lewis	4	4	Public Speaking	Data Analysis	Artistic	Leadership	Problem Solving	Houston	...	Unf

501 rows x 26 columns

```
In [15]: # Step 4: Define predictor variables (X) and target variable (y)
X = engaged_students[['StudentID', 'GPA', 'CompositeScore']]
y = engaged_students['Engagement']
```

```
In [16]: # Step 5: Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [17]: # Step 6: Train the Random Forest Classifier
rf_classifier = RandomForestClassifier()
rf_classifier.fit(X_train, y_train)
```

```
Out[17]: ▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [18]: # Step 7: Make predictions
y_pred = rf_classifier.predict(X_test)
```

```
In [19]: # Step 8: Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Accuracy: 0.26732673267326734

Classification Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	28
1	0.00	0.00	0.00	2
2	0.35	0.38	0.36	37
3	0.27	0.38	0.31	34
accuracy			0.27	101
macro avg	0.15	0.19	0.17	101
weighted avg	0.22	0.27	0.24	101

```
/Users/amjadalikudsi/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/Users/amjadalikudsi/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/Users/amjadalikudsi/anaconda3/lib/python3.11/site-packages/sklearn/metrics/_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use 'zero_division' parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

Bonus Task: Prediction that nextGPA is not being hampered based on Engagement

```
In [20]: # Step 5: Predict whether their next GPA will not be hampered
# Here, you can use additional features such as GPA, attendance, composite score
# Define predictor variables (X) and target variable (y)
X_gpa = engaged_students[['StudentID', 'GPA', 'CompositeScore']]
# Assuming you have the target variable 'NextGPA' which indicates whether the next GPA is hampered or not

# Make predictions
next_gpa_predictions = rf_classifier.predict(X_gpa)

# Add predictions as a new column to the DataFrame
engaged_students['EngagementMetrics'] = next_gpa_predictions

# Display the DataFrame with predictions
print(engaged_students[['StudentID', 'GPA', 'NextGPA', 'EngagementMetrics']])
```

	StudentID	GPA	NextGPA	EngagementMetrics
0	1	3.27	3.289579	0
1	2	3.17	3.118852	0
6	7	3.35	3.356055	0
7	8	3.64	3.673059	2
9	10	3.03	2.971565	2
...
990	991	3.37	3.354229	3
993	994	3.23	3.181503	3
996	997	3.62	3.685658	3
997	998	3.09	2.999678	0
998	999	3.17	3.116268	3

[501 rows x 4 columns]

/var/folders/33/03_543s9719fh3y8mrxhst580000gn/T/ipykernel_46941/2134801628.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
engaged_students['EngagementMetrics'] = next_gpa_predictions

```
In [21]: # Assuming 'new_students' DataFrame contains the data including the predicted next GPA and the current GPA

# Define a function to check if next GPA is less than current GPA
def check_pass_fail(row):
    if row['NextGPA'] < threshold:
        return 0 # If next GPA is less than current GPA, return 0 (Fail)
    else:
        return 1 # If next GPA is greater than or equal to current GPA, return 1 (Pass)

# Apply the function to create the new_pass_fail column
engaged_students['new_pass_fail'] = engaged_students.apply(check_pass_fail, axis=1)

# Display the updated DataFrame with the new_pass_fail column
print(engaged_students[['StudentID', 'GPA', 'NextGPA', 'EngagementMetrics', 'new_pass_fail']])
```

	StudentID	GPA	NextGPA	EngagementMetrics	new_pass_fail
0	1	3.27	3.289579	0	1
1	2	3.17	3.118852	0	1
6	7	3.35	3.356055	0	1
7	8	3.64	3.673059	2	1
9	10	3.03	2.971565	2	0
...
990	991	3.37	3.354229	3	1
993	994	3.23	3.181503	3	1
996	997	3.62	3.685658	3	1
997	998	3.09	2.999678	0	0
998	999	3.17	3.116268	3	1

[501 rows x 5 columns]

/var/folders/33/03_543s9719fh3y8mrxhst580000gn/T/ipykernel_46941/3552749785.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)
engaged_students['new_pass_fail'] = engaged_students.apply(check_pass_fail, axis=1)

Visualizations

```
In [22]: import matplotlib.pyplot as plt
import seaborn as sns

# Visualization 1: Distribution of GPA
plt.figure(figsize=(8, 6))
sns.histplot(data=engaged_students, x='GPA', bins=20, kde=True, color='skyblue')
plt.title('Distribution of GPA')
plt.xlabel('GPA')
plt.ylabel('Frequency')
plt.show()

# Visualization 2: Distribution of Next GPA
plt.figure(figsize=(8, 6))
sns.histplot(data=engaged_students, x='NextGPA', bins=20, kde=True, color='salmon')
plt.title('Distribution of Next GPA')
plt.xlabel('Next GPA')
plt.ylabel('Frequency')
plt.show()

# Visualization 3: Scatter plot of GPA vs. Next GPA
plt.figure(figsize=(8, 6))
sns.scatterplot(data=engaged_students, x='GPA', y='NextGPA', color='green')
plt.title('Scatter plot of GPA vs. Next GPA')
plt.xlabel('GPA')
plt.ylabel('Next GPA')
plt.show()

# Visualization 4: Bar plot of Pass/Fail based on new_pass_fail
plt.figure(figsize=(6, 4))
sns.countplot(data=engaged_students, x='new_pass_fail', palette='Set2')
plt.title('Pass/Fail based on new_pass_fail')
plt.xlabel('Pass/Fail')
plt.ylabel('Count')
plt.xticks(ticks=[0, 1], labels=['Fail', 'Pass'])
plt.show()

# Visualization 5: Count of Engagements
plt.figure(figsize=(8, 6))
sns.countplot(data=engaged_students, x='Engagement', palette='pastel')
plt.title('Count of Engagements')
plt.xlabel('Engagement')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()
```





